

Neural Network Basics

1. What does a neuron compute?

- ☐ A neuron computes an activation function followed by a linear function ($z = Wx + b$).
 - ☒ A neuron computes a linear function ($z = Wx + b$) followed by an activation function.

Correct, we generally say that the output of a neuron is $a = g(Wx + b)$ where g is the activation function (sigmoid, tanh, ReLU, ...).
 - ☐ A neuron computes a function g that scales the input x linearly ($Wx + b$).
 - ☐ A neuron computes the mean of all features before applying the output to an activation function.
-

2. Which of these is the "Logistic Loss"?

- ☐ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = \max(0, y^{(i)} - \hat{y}^{(i)})$
 - ☒ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = -(y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$

Correct, this is the logistic loss you've seen in lecture!
 - ☐ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = |y^{(i)} - \hat{y}^{(i)}|$
 - ☐ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = |y^{(i)} - \hat{y}^{(i)}|^2$
-

3. Suppose `img` is a (32,32,3) array, representing a 32x32 image with 3 color channels red, green and blue. How do you reshape this into a column vector?

- ☒ `x = img.reshape((32 * 32 * 3, 1))`
 - ☐ `x = img.reshape((1, 32 * 32 * 3))`
 - ☐ `x = img.reshape((3, 32 * 32))`
 - ☐ `x = img.reshape((32 * 32, 3))`
-

4. Consider the two following random arrays "a" and "b":

```
1 a = np.random.randn(2, 3) # a.shape = (2, 3)
2 b = np.random.randn(2, 1) # b.shape = (2, 1)
3 c = a + b
```

What will be the shape of "c"?

- ☒ `c.shape = (2, 3)`

Yes! This is broadcasting. `b` (column vector) is copied 3 times so that it can be summed to each column of `a`.
- ☐ The computation cannot happen because the sizes don't match. It's going to be "Error"!
- ☐ `c.shape = (2, 1)`

☐ c.shape = (3, 2)

5. Consider the two following random arrays "a" and "b":

```
1 a = np.random.randn(4, 3) # a.shape = (4, 3)
2 b = np.random.randn(3, 2) # b.shape = (3, 2)
3 c = a * b
```

What will be the shape of "c"?

☐ c.shape = (3, 3)

☒ The computation cannot happen because the sizes don't match. It's going to be "Error"!

Indeed! In numpy the "*" operator indicates element-wise multiplication. It is different from "np.dot()". If you would try "c=np.dot(a, b)" you would get c.shape = (4, 2).

☐ c.shape = (4, 2)

☐ c.shape = (4, 3)

6. Suppose you have n_x input features per example. Recall that $X = [x^{(1)}, x^{(2)} \dots x^{(m)}]$. What is the dimension of X ?

☐ $(1, m)$

☐ (m, n_x)

☐ $(m, 1)$

☒ (n_x, m)

7. Recall that `np.dot(a, b)` performs a matrix multiplication on a and b, whereas `a*b` performs an element-wise multiplication. Consider the two following random arrays "a" and "b":

```
1 a = np.random.randn(12288, 150) # a.shape = (12288, 150)
2 b = np.random.randn(150, 45) # b.shape = (150, 45)
3 c = np.dot(a, b)
```

What is the shape of c?

☐ c.shape = (12288, 150)

☐ c.shape = (150, 150)

☒ c.shape = (12288, 45)

Correct, remember that a `np.dot(a, b)` has shape (number of rows of a, number of columns of b). The sizes match because:

"number of columns of a = 150 = number of rows of b"

☐ The computation cannot happen because the sizes don't match. It's going to be "Error"!

8. Consider the following code snippet:

```

1 # a.shape = (3,4)
2 # b.shape = (4,1)
3 for i in range(3):
4     for j in range(4):
5         c[i][j] = a[i][j] + b[j]

```

How do you vectorize this?

- ☐ $c = a.T + b.T$
- ☐ $c = a.T + b$
- ☐ $c = a + b$
- ☒ $c = a + b.T$

9. Consider the following code:

```


1 a = np.random.randn(3, 3)
2 b = np.random.randn(3, 1)
3 c = a * b

```

What will be c?

- ☒ This will invoke broadcasting, so b is copied three times to become (3, 3), and * is an element-wise product so c.shape will be (3, 3).
- ☐ This will invoke broadcasting, so b is copied three times to become (3, 3), and * invokes a matrix multiplication operation of two 3x3 matrices so c.shape will be (3, 3).
- ☐ This will multiply a 3x3 matrix a with a 3x1 vector, thus resulting in a 3x1 vector. That is, c.shape = (3, 3).
- ☐ It will lead to an error since you cannot use "*" to operate on these two matrices. You need to instead use np.dot(a, b).

10. Consider the following computation graph.



- ☐ $J = (c - 1) * (b + a)$
- ☒ $J = (a - 1) * (b + c)$

Yes. $J = u + v - w = a * b + a * c - (b + c) = a * (b + c) - (b + c) = (a - 1) * (b + c)$

☐ $J = a * b + b * c + a * c$

☐ $J = (b - 1) * (c + a)$
